



Figure 1: Example sketches produced collaboratively by humans and artificial neural network model. Red strokes produced by humans; yellow strokes produced by model.

collabdraw: an environment for collaborative sketching with an artificial agent

Judith Fan
UC San Diego
La Jolla, CA
jefan@ucsd.edu

Monica Dinculescu
Google Brain
Mountain View, CA
noms@google.com

David Ha
Google Brain
Mountain View, CA
hadavid@google.com

ABSTRACT

Sketching is one of the most accessible techniques for communicating our ideas quickly and for collaborating in real time to produce meaningful content. Here we present a web-based environment for collaborative sketching of everyday visual concepts. We explore the integration of an artificial agent, instantiated as a recurrent neural network, who is both cooperative and responsive to actions performed by its human collaborator. To evaluate the quality of the sketches produced in this environment, we conducted an experimental user study and found that sketches produced collaboratively were just as recognizable as those produced by humans on their own. Further control

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

ACM Creativity & Cognition '19, June 23–26, 2019, San Diego, CA

© 2019 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

analyses suggest that the semantic properties of these sketches were indeed the product of collaboration, rather than attributable to the contributions of the human or the artificial agent alone. Taken together, our findings attest to the potential of systems enabling real-time collaboration between humans and machines to create novel and meaningful content.

CCS CONCEPTS

• **Human-centered computing** → **Empirical studies in interaction design.**

KEYWORDS

drawing, human-computer interaction, artificial intelligence

ACM Reference Format:

Judith Fan, Monica Dinculescu, and David Ha. 2019. collabdraw: an environment for collaborative sketching with an artificial agent. In *ACM Creativity & Cognition '19, San Diego, CA*. ACM, New York, NY, USA, 6 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

INTRODUCTION

The ability to express our ideas in visual form provides the foundation for important creative activities, including art, design, and science. A critical advantage of having such visual formats is that they are easily shared and naturally support collaboration. Sketching is one of the most accessible techniques for communicating our ideas quickly and collaborating in real time to produce meaningful content. What are the computational ingredients that enable successful coordination between agents to construct meaningful sketches together, even in the absence of auxiliary communication channels (e.g., spoken language)? To investigate this question, we developed a web-based environment for collaborative sketching of everyday visual concepts. We integrated an artificial agent into this environment, instantiated as a recurrent neural network model [3, 7], to collaborate with each person who entered the environment. We explore the hypothesis that endowing this agent with: *first*, a well-defined goal it shares with its human collaborator, and *second*, the ability to adapt to this person's contributions, would support the collaborative production of meaningful sketches. To test this hypothesis, we conducted an experimental user study and found that the resulting collaborative sketches contained similar semantic information to those produced by humans on their own. Further control analyses suggest that the semantic information in these sketches were indeed the product of collaboration, rather than attributable to the contributions of the human or the artificial agent alone. Taken together, our findings attest to the potential of systems enabling real-time collaboration between humans and machines to create novel and meaningful content [5?].

collab sketching




Figure 2: Concepts were cued with their label (e.g., ‘bear’). All participants drew each concept under both collab and solo conditions. On collab trials, the human and model alternated producing a single stroke until the human was satisfied with the drawing. On solo trials, the human drew the cued concept on their own.

COLLABDRAW: A COLLABORATIVE SKETCHING ENVIRONMENT

Building this web-based collaborative sketching environment posed two key technical challenges: *first*, development of a cooperative artificial agent capable of sketching familiar visual concepts, and *second*, real-time deployment in a web browser. To solve the first challenge, we leverage a recently developed recurrent neural network model (sketch-rnn, [3]), that was previously trained on a large number of human-generated sketches of various visual concepts [4], and shown to have learned a latent distribution that can be sampled to generate novel sketches of those concepts. Importantly for our study, sketch-rnn can propose coherent continuations of partial sketches of a particular target concept, providing a natural interface for collaboration with a human who shares the goal of sketching that concept and has generated a partial sketch of it. To solve the second challenge, we used an implementation of sketch-rnn from Magenta.js [7], a high-level JavaScript API for doing inference with machine learning models in client-side applications, built with Tensorflow.js [10].

EXPERIMENTAL USER STUDY

We conducted an experimental user study in order to evaluate how well our collaborative sketching environment supported the production of recognizable sketches. 90 participants recruited via Amazon Mechanical Turk (AMT) completed the study. In each session, participants drew eight different visual concepts (i.e., bear, cat, duck, lion, pig, rabbit, sheep, and swan) four times each, in a randomly interleaved sequence. In order to measure the impact of collaboration, controlling for variation between individuals in sketching behavior, participants sketched each object twice under collaborative (collab) and twice under non-collaborative (solo) settings. On each collab sketching trial, the human participant and artificial agent (i.e., ‘model’) took turns producing one stroke each until the human participant was satisfied with the sketch (Fig. 2). The human participant always produced the first stroke, and thus all odd-numbered strokes in each collaborative sketch, whereas the model produced all of the even-numbered strokes. On each solo sketching trial, the human participant drew the object entirely on their own, providing a measure of the baseline recognizability of human sketches of each object in our dataset.

ANALYSIS OF SKETCH BEHAVIOR

The primary goal of the user study was to test two critical aspects of the collabdraw environment:

- That the environment supports real-time collaboration between a human and artificial agent to produce meaningful sketches.
- That the semantic information conveyed in these collaborative sketches was indeed the product of collaboration, rather than attributable to the contributions of the human or the artificial agent alone.

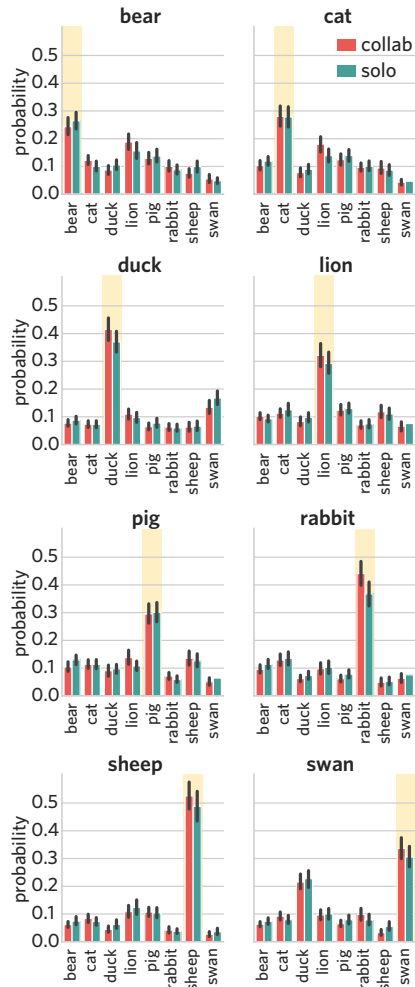


Figure 3: Distribution of probabilities assigned to each concept by the classifier, for each target concept and condition. Error bars reflect 95% confidence intervals.

Extracting semantic information from sketches

Measuring the semantic properties of a sketch that determines its recognizability requires a principled approach for encoding their high-level visual properties. Here we capitalize on recent work validating the use of deep convolutional neural network (CNN) models, pre-trained on challenging visual tasks, to encode such properties in sketches [2, 8, 11]. Specifically, we used VGG-19 [9] pre-trained to categorize objects in photos from the Imagenet database [1] to extract high-level feature-vector representations of each sketch. Prior to feature extraction, we rendered all sketches as 224px x 224px RGB images containing only black strokes of fixed width (5px), and that were cropped to include only the sketch plus 10px padding. Each 4096-dimensional feature vector reflected VGG-19 activations to sketches in the second fully-connected layer of the network (i.e., fc6). To extract the degree to which each sketch expressed the target concept, we applied an 8-way logistic classifier with L2 regularization to predict the identity of the sketched concept. The classifier was trained on a category-balanced set of 40K sketches of the same 8 concepts in the Quickdraw dataset. Critically, although the artificial agent and the CNN-based classifier were both trained on subsets of the Quickdraw dataset, because they encode sketches in different ways (i.e., strokes vs. pixels), underwent different training regimes, and did not interact at test time, the CNN-based classifier provides a distinct measure of the semantic information in each sketch.

Does collaboration yield meaningful sketches?

To evaluate the degree to which our system supported successful collaboration, for each collab sketch we extracted the pattern of probabilities assigned to each of the 8 concepts by the classifier, and compared this pattern to that assigned to sketches in the baseline solo condition (Fig. 3). We found that the highest probability value was consistently assigned to the target concept in both conditions (collab: 0.357, 95% CI: [0.340, 0.374]; solo: 0.333, 95% CI: [0.316, 0.350]), validating our general approach to measuring semantic information in these sketches. If anything, there was modest evidence that collab sketches were slightly more recognizable than solo sketches, as revealed by a generalized linear mixed-effects model predicting classifier accuracy from condition, that included random intercepts for participant and concept ($b = 0.142, z = 1.78, p = 0.075$). Beyond classifier accuracy, we further found a strikingly high degree of correspondence in the *pattern* of probabilities assigned to each concept between the collab and solo conditions. We quantified this correspondence using Jensen-Shannon divergence (JSD), which captures how different two probability distributions are, and lies in the range [0, 1]. We compared the observed mean divergence across concepts (JSD= 0.0024) to a null distribution where the assignment of probabilities to concept was scrambled independently for each condition. We found that the observed JSD fell far below the null distribution for every concept, providing strong evidence of non-trivial and similar structure in




Figure 4: Average probability assigned to target concept by classifier for intact collab and solo sketches (left). Comparison with lesioned collab sketches in which only the human-generated (middle) or model-generated (right) strokes remain, as well as lesioned solo sketches in which only odd-numbered (middle) or even-numbered (right) strokes remain. Error bars reflect 95% confidence intervals.

each pair of distributions. Moreover, we found that similar amounts of effort were invested in producing sketches in both conditions: neither the number of strokes ($collab = 12.2$; $solo = 11.7$; 95% CI for difference = $[-5.63, 5.06]$), nor the amount of time taken to finish each sketch reliably differed between conditions ($collab = 39.0s$; $solo = 36.8s$; 95% CI for difference = $[-41.6s, 38.9s]$). Taken together, these results suggest that, although the human participant and artificial agent had no other way of communicating apart from their contributions to each sketch, they nevertheless succeeded in producing sketches that carried as much relevant semantic information about each concept as those sketches produced by the human participant alone.

Did both parties contribute to the result?

An alternative explanation for the findings so far is that the human participant was primarily responsible for contributing strokes that conveyed relevant semantic information about the target concept. On this account, if we were to consider only the human-generated strokes in collab sketches, then the probability assigned to the target concept should remain unaffected. Conversely, if the artificial agent were primarily responsible for contributing semantically relevant information in each sketch, then considering only the model-generated strokes should provide as much information about the target as intact collab sketches. To test both of these alternative explanations, we conducted the same classifier-based analysis as above on ‘lesioned’ versions of each collab sketch in which either all of the human (odd-numbered) strokes or all of the model (even-numbered) strokes were removed. We found that the probability assigned to the target decreased substantially in both types of lesioned sketches, and to a comparable degree (human strokes only: 0.277, 95% CI: $[0.264, 0.291]$; model strokes only: 0.269, 95% CI: $[0.257, 0.281]$), suggesting that both parties contributed to the overall meaning of each sketch (Fig. 4). To provide a baseline for comparison, we additionally repeated the analysis on the solo sketches, lesioning either all of the odd-numbered or even-numbered strokes. We found that these lesions produced a comparable decrement in the amount of semantic information in each sketch (odd strokes only: 0.267, 95% CI: $[0.254, 0.280]$; even strokes only: 0.208, 95% CI: $[0.199, 0.218]$); if anything, the sketches containing only the human-generated even strokes conveyed less information about the target than those containing only the model-generated even strokes. In sum, these results suggest that the semantic properties of the collab sketches were indeed the product of collaboration, rather than attributable to the contributions of the human or the artificial agent alone.

FUTURE WORK

The current study affirms the potential of systems enabling real-time collaboration between humans and machines to create novel and meaningful content. Next, we plan to collect human judgments to validate our classifier-based analyses of the semantic information in collaborative sketches. More

broadly, we aim to further develop our collaborative sketching environment for a variety of other applications, including ones that promote creative improvisation during sketch production [5?], and others that provide scaffolding for individuals who may experience difficulty sketching on their own, such as young children [6] and neurological patients [12]). In the long run, such technologies may even help users to discover new modes of creative visual expression.

ACKNOWLEDGMENTS

Thanks to Bria Long and Mike Frank for helpful discussion.

REFERENCES

- [1] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. Imagenet: A large-scale hierarchical image database. (2009).
- [2] Judith E. Fan, Daniel L. K. Yamins, and Nicholas B. Turk-Browne. 2018. Common Object Representations for Visual Production and Recognition. *Cognitive Science* (2018).
- [3] David Ha and Douglas Eck. 2017. A neural representation of sketch drawings. *arXiv preprint arXiv:1704.03477* (2017).
- [4] Jonas Jongejan, Henry Rowley, Takashi Kawashima, Jongmin Kim, and Nick Fox-Gieg. 2017. Google QuickDraw. <https://quickdraw.withgoogle.com/>
- [5] Pegah Karimi, Nicholas Davis, Kazjon Grace, and Mary Lou Maher. 2018. Deep Learning for Identifying Potential Conceptual Shifts for Co-creative Drawing. *CoRR* abs/1801.00723 (2018). arXiv:1801.00723 <http://arxiv.org/abs/1801.00723>
- [6] Bria Long, Judith E Fan, and Michael C Frank. 2018. Drawings as a window into developmental changes in object representations. In *Proceedings of the 40th Annual Meeting of the Cognitive Science Society*.
- [7] Adam Roberts, Curtis Hawthorne, and Ian Simon. 2018. Magenta.js: A javascript api for augmenting creativity with deep learning. (2018).
- [8] Patsorn Sangkloy, Nathan Burnell, Cusuh Ham, and James Hays. 2016. The sketchy database: learning to retrieve badly drawn bunnies. *ACM Transactions on Graphics (TOG)* 35, 4 (2016), 119.
- [9] Karen Simonyan and Andrew Zisserman. 2014. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556* (2014).
- [10] Daniel Smilkov, Nikhil Thorat, Yannick Assogba, Ann Yuan, Nick Kreeger, Ping Yu, Kangyi Zhang, Shanqing Cai, Eric Nielsen, David Soergel, et al. 2019. TensorFlow.js: Machine Learning for the Web and Beyond. *arXiv preprint arXiv:1901.05350* (2019).
- [11] Qian Yu, Yongxin Yang, Feng Liu, Yi-Zhe Song, Tao Xiang, and Timothy M Hospedales. 2017. Sketch-a-net: A deep neural network that beats humans. *International journal of computer vision* 122, 3 (2017), 411–425.
- [12] Daniel Zalewski. 2015. Life Lines. *The New Yorker* (2015).